

Suppose we want to minimize  $J(\theta)$  as a function of  $\theta$ . For this example, suppose  $J : \mathfrak{R} \mapsto \mathfrak{R}$ , so that  $\theta \in \mathfrak{R}$ . If we are using `minFunc` or some other optimization algorithm, then we usually have implemented some function  $g(\theta)$  that purportedly computes  $\frac{d}{d\theta}J(\theta)$ .

How can we check if our implementation of  $g$  is correct?

Recall the mathematical definition of the derivative as:

$$\frac{d}{d\theta}J(\theta) = \lim_{\epsilon \rightarrow 0} \frac{J(\theta + \epsilon) - J(\theta - \epsilon)}{2\epsilon}.$$

Thus, at any specific value of  $\theta$ , we can numerically approximate the derivative as follows:

$$\frac{J(\theta + \text{EPSILON}) - J(\theta - \text{EPSILON})}{2 \times \text{EPSILON}}$$

In practice, we set `EPSILON` to a small constant, say around  $10^{-4}$ . (There's a large range of values of `EPSILON` that should work well, but we don't set `EPSILON` to be "extremely" small, say  $10^{-20}$ , as that would lead to numerical roundoff errors.)

Thus, given a function  $g(\theta)$  that is supposedly computing  $\frac{d}{d\theta}J(\theta)$ , we can now numerically verify its correctness by checking that

$$g(\theta) \approx \frac{J(\theta + \text{EPSILON}) - J(\theta - \text{EPSILON})}{2 \times \text{EPSILON}}.$$

The degree to which these two values should approximate each other will depend on the details of  $J$ . But assuming `EPSILON` =  $10^{-4}$ , you'll usually find that the left- and right-hand sides of the above will agree to at least 4 significant digits (and often many more).

Now, consider the case where  $\theta \in \mathfrak{R}^n$  is a vector rather than a single real number (so that we have  $n$  parameters that we want to learn), and  $J : \mathfrak{R}^n \mapsto \mathfrak{R}$ . We now generalize our derivative checking procedure to the case where  $\theta$  may be a vector (as in our linear regression and logistic regression examples). If ever we are optimizing over several variables or over matrices, we can always pack these parameters into a long vector and use the same method here to check our derivatives. (This will often need to be done anyway if you want to use off-the-shelf optimization packages.)

Suppose we have a function  $g_i(\theta)$  that purportedly computes  $\frac{\partial}{\partial \theta_i}J(\theta)$ ; we'd like to check if  $g_i$  is outputting correct derivative values. Let  $\theta^{(i+)} = \theta + \text{EPSILON} \times e_i$ , where

$$g_i(\theta) \approx \frac{J(\theta^{(i+)}) - J(\theta^{(i-)})}{2 \times \text{EPSILON}}.$$

## Gradient checker code

As an exercise, try implementing the above method to check the gradient of your linear regression and logistic regression functions. Alternatively, you can use the provided `ex1/grad_check.m` file (which takes arguments similar to `minFunc`) and will check  $\frac{\partial J(\theta)}{\partial \theta_i}$  for many random choices of  $i$ .

